

1. Architecture and training

1.1 Detailed network architecture

The t-VGG GAP network used in the present study consisted of three blocks (each containing two convolutional and one max-pooling layer), one Global Averaging Pooling layer, followed by one output layer. The output layer consisted of a single neuron (with sigmoid activation function) when predicting the outcome; it consisted of four neurons (with softmax activation function) when predicting the etiology.

Table S1: Blocks and GAP

	Layer	Description
Block 1	Input	1x500x9, corresponding to 10 seconds of EEG at 50 Hz, with nine channels
	1-D Convolution	16 filters, kernel length 3, stride 1
	BatchNormalization	
	Activation	ReLU
	1-D Convolution	16 filters, kernel length 3, stride 1
	BatchNormalization	
	Activation	ReLU
	Max-Pooling	pool size 4, stride 4
Block 2	1-D Convolution	32 filters, kernel length 3, stride 1
	BatchNormalization	
	Activation	ReLU
	1-D Convolution	32 filters, kernel length 3, stride 1
	BatchNormalization	
	Activation	ReLU
Block 3	Max-Pooling	pool size 4, stride 4
	1-D Convolution	32 filters, kernel length 3, stride 1
	BatchNormalization	
	Activation	ReLU
	1-D Convolution	32 filters, kernel length 3, stride 1
	BatchNormalization	
	Activation	ReLU activation
GAP	Max-Pooling	pool size 4, stride 4
	Global Average Pooling	
Output Layer		
Setting:		
Prognostication	Fully-connected Layer	1 neuron, sigmoid activation
OR		
Setting:		
Etiology	Fully-connected Layer	4 neurons, softmax activation

1.2. Training procedure

We performed a stratified 5-fold cross validation, each time using four folds for training and the remaining fold as test set. The stratified split ensures that each fold contains roughly the same class distribution. In the binary setting (outcome/survival) the loss function to be minimized by the optimizer was the binary cross-entropy; in the multi-

class setting (etiology), the loss function was the sparse categorical cross-entropy. We optimized our model using the Adam optimizer with a learning rate of $2e-5$ (binary) or $1e-4$ (multi-class). To counter-act the class imbalances, we also calculated the class weights of the training set, which were given as parameter to the Keras model fitting call. This parameter setting tells Keras to weigh the loss function relative to the class occurrences. The number of trainable parameters was 12'305 in the binary setting and 12'404 in the etiology setting. See source code below.

2. Source Code

2.1 t-VGG GAP Model Implementation

'''

This method returns the compiled "tVGG GAP" model used in the binary class setting (prognostication).
It is a 1D-convolutional network for EEG epoch inputs of the shape (timeseries_windowsize, number of channels)

The model architecture is adapted from the paper "EEG-based Outcome Prediction after Cardiac Arrest with Convolutional Neural Networks: Performance and Visualization of Discriminative Features" by Jonas et al. Human Brain Mapping 2019 where further information can be found.

Keras v2.3.1 with a TensorFlow backend

'''

```
def compile_single_class_TVGG_model(input_vars=(500,9), learningrate = 0.0001):
    timeseries_windowsize = input_vars[0]
    num_channels = input_vars[1]
    model = Sequential()

    # Block 1
    model.add(Conv1D(16, kernel_size=3, strides=1,
                     name="firstCV", use_bias=False,
                     kernel_initializer="glorot_uniform", input_shape=(timeseries_windowsize, num_channels)))
    model.add(BatchNormalization())
    model.add(Activation('relu'))
    model.add(Conv1D(16, kernel_size=3, strides=1, use_bias=False, kernel_initializer="glorot_uniform"))
    model.add(BatchNormalization())
    model.add(Activation('relu'))
    model.add(MaxPooling1D(pool_size = 4, strides=4))

    # Block 2
    model.add(Conv1D(32, kernel_size=3, kernel_initializer="glorot_uniform",
                     strides=1, use_bias = False))
    model.add(BatchNormalization())
    model.add(Activation('relu'))
    model.add(Conv1D(32, kernel_size=3, strides=1, use_bias = False))
    model.add(BatchNormalization())
    model.add(Activation('relu'))
    model.add(MaxPooling1D(pool_size = 4, strides=4))

    # Block 3
    model.add(Conv1D(32, kernel_size=3, kernel_initializer="glorot_uniform", strides=1, use_bias = False))
    model.add(BatchNormalization())
    model.add(Activation('relu'))
    model.add(Conv1D(32, kernel_size=3, strides=1, use_bias=False, name="lastCV")) # the name is used as
referral for GradCAM
    model.add(BatchNormalization())
    model.add(Activation('relu'))
    model.add(MaxPooling1D(pool_size = 4, strides=4))

    # Global Average Pooling operation after the third block
    model.add(GlobalAveragePooling1D())

    # into a final single-neuron FC layer outputting the probability
    model.add(Dense(1, activation='sigmoid'))
```

```

adam = optimizers.Adam(lr=learningrate)
model.compile(loss='binary_crossentropy',
              optimizer=adam,
              metrics=['accuracy'])
return model

'''
Returns the same model architecture compiled and adjusted for the multi-class (etiology, 4 classes) setting.
'''
def compile_multi_class_TVGG_model(input_vars=(500,9), learningrate = 0.0001, num_final_classes = 4):
    timeseries_windowsize = input_vars[0]
    num_channels = input_vars[1]
    model = Sequential()
    # Block 1
    model.add(Conv1D(16, kernel_size=3, strides=1,
                    name="firstCV", use_bias=False,
                    kernel_initializer="glorot_uniform", input_shape=(timeseries_windowsize, num_channels)))
    model.add(BatchNormalization())
    model.add(Activation('relu'))
    model.add(Conv1D(16, kernel_size=3, strides=1, use_bias=False, kernel_initializer="glorot_uniform"))
    model.add(BatchNormalization())
    model.add(Activation('relu'))
    model.add(MaxPooling1D(pool_size = 4, strides=4))

    # Block 2
    model.add(Conv1D(32, kernel_size=3, kernel_initializer="glorot_uniform",
                    strides=1, use_bias = False))
    model.add(BatchNormalization())
    model.add(Activation('relu'))
    model.add(Conv1D(32, kernel_size=3, strides=1, use_bias = False))
    model.add(BatchNormalization())
    model.add(Activation('relu'))
    model.add(MaxPooling1D(pool_size = 4, strides=4))

    # Block 3
    model.add(Conv1D(32, kernel_size=3, kernel_initializer="glorot_uniform", strides=1, use_bias = False))
    model.add(BatchNormalization())
    model.add(Activation('relu'))
    model.add(Conv1D(32, kernel_size=3, strides=1, use_bias=False, name="lastCV"))
    model.add(BatchNormalization())
    model.add(Activation('relu'))
    model.add(MaxPooling1D(pool_size = 4, strides=4))

    # GAP into a single output layer
    model.add(GlobalAveragePooling1D())
    # num_classes neurons activated by softmax
    model.add(Dense(num_final_classes, activation='softmax'))

    adam = optimizers.Adam(lr=learningrate)
    model.compile(loss='sparse_categorical_crossentropy',
                  optimizer=adam,
                  metrics=['accuracy'])
    return model

```

2.2 GradCAM Implementation

'''

The GradCAM method implementations.
Python v3.7.1

Keras v2.3.1 with a TensorFlow backend

'''

```
from keras import backend as K
from tensorflow.keras import Model
```

'''

This method returns the GradCAM values for a selected class in the binary setting.

Inputs:

model = The fully trained and saved model for classification
 EEG_epoch = The specific 1D epoch to obtain the values for
 timeseries_windowsize = length in datapoints of the epoch
 numChannels = # of channels in the input epoch
 selected_class = Obtain values for class 0 or class 1 (binary)
 layer_name = The layer name in the model (pre-defined) from which the gradients will be obtained (usually the last)

Returns:

gc_values = An array with the gradcam values which can then be further resized and edited for visualizations.

In our case the EEG epoch is of shape (500, 9), the last convolutional layer is named "lastCV" and the selected class is 1.

'''

```
def get_gradcam_singleclass(model, EEG_epoch, timeseries_windowsize, numChannels, selected_class=1,
layer_name="lastCV"):
    y_c = model.output[0]
    conv_layer = model.get_layer(layer_name) # defined in the saved model
    gradcam_model = Model([model.inputs], [conv_layer.output, model.output])

    # Get gradient w.r.t. the output of the selected conv layer
    with tf.GradientTape() as gtape:
        conv_output, predictions = gradcam_model(EEG_epoch.reshape(1, timeseries_windowsize, numChannels))
        loss = predictions[0]
        grads = gtape.gradient(loss, conv_output)
        pooled_grads = K.mean(grads, axis=(0, 1))

    if selected_class == 0:
        pooled_grads = pooled_grads * -1

    gc_values = tf.reduce_mean(tf.multiply(pooled_grads, conv_output), axis=-1)
    gc_values = gc_values[0]
    gc_values = np.maximum(gc_values, 0)
    gc_values /= np.max(gc_values) # optional norm.

    return gc_values
```

'''

This method was used to retrieve the gradCAM values in the multi-class setting, i.e. the etiology setting.

Changed inputs:

selected_class = Possibilities of 0,1,2 or 3 in the 4-class setting.

'''

```
def get_gradcam_multiclass(model, EEG_epoch, timeseries_windowsize, numChannels, selected_class=2,
layer_name="lastCV"):
    conv_layer = model.get_layer(layer_name)
    gradcam_model = Model([model.inputs], [conv_layer.output, model.output])
    with tf.GradientTape() as gtape:
        conv_output, predictions = gradcam_model(EEG_epoch.reshape(1, timeseries_windowsize, numChannels))
        loss = predictions[:,selected_class]
        grads = gtape.gradient(loss, conv_output)
        pooled_grads = K.mean(grads, axis=(0, 1))

    gc_values = tf.reduce_mean(tf.multiply(pooled_grads, conv_output), axis=-1)
    gc_values = gc_values[0]
    gc_values = np.maximum(gc_values, 0)
```

```
gc_values /= np.max(gc_values) # optional norm.
return gc_values
```

2.3 Training Procedure

```
import numpy as np
from sklearn.utils import class_weight
from sklearn.model_selection import StratifiedKFold
'''
```

This snippet showcases the training procedure applied in our study. Dataset-specific processing steps are omitted. Python v3.7.1 and Keras v2.3.1 with a TensorFlow backend
sklearn v0.20.3 ; numpy v1.18.5

```
'''
# ... pre-processing, data-splitting and loading:
EEG_DATALIST = # a list of patient-ecg names
EEG_LABELLIST = # a list containing the corresponding label to each ecg
```

```
# 5-fold stratified cross-Validation:
# (omitted: prepare documentation for each fold)
skf = StratifiedKFold(n_splits=5, random_state=CV_RANDOM_STATE, shuffle=True)
```

```
for train_index, test_index in skf.split(EEG_DATALIST, EEG_LABELLIST):
    # create training sets (4 folds) and a test set (1 CV fold)
    train_ecg_list, test_ecg_list = EEG_DATALIST[train_index], EEG_DATALIST[test_index]
    x_train, y_train, ids_train = # create_dataset(train_ecg_list, ...) - load the epochs
    x_test, y_test, ids_test = # create_dataset(test_ecg_list, ...)

    # load a compiled "t-VGG GAP" model. Here: single-class model, see further supplied material.
    lr = 0.00002
    model = eeg_models.compile_single_class_TVGG_model(input_vars=(500,9), learningrate = lr)

    # sidestep: counter-act class imbalances, supply class weights to training call (see model.fit)
    class_weights = class_weight.compute_class_weight('balanced',np.unique(y_train),y_train)
    class_weight_dict = {}
    class_weight_dict[0], class_weight_dict[1] = class_weights[0], class_weights[1]
```

```
# fit and train the model
# the model was pre-compiled with the parameters:
# Optimizer: Adam; loss function: binary cross-entropy
callback = modelhelper.get_earlytraining_callback(training_stop = ts)
model.fit(x_train, y_train, batch_size=bs,
          epochs=max_epochs, verbose=1,shuffle=True,
          callbacks=callback, class_weight=class_weight_dict)
```

```
# omitted: documentation and saving the model (a saved state is needed for Grad-CAM)
# model.save(model_name)
```

```
# model evaluation
acc, sensitivity, specificity, PPV, NPV, roc_auc = # evaluate_model(test_ecg_list, model=model, ...)
```

```
# omitted: further documentation
# reset the loaded and compiled model:
K.clear_session()
model = None
```